

Multi Agent Systems

- Lab 5 -

Q-Learning and SARSA

Value Iteration Recap

- **Policy** and **Value Iteration** algorithms require knowledge of **environment dynamics**

$$V^*(s) = \max_a [R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')]$$

- $P(s' / s, a)$ has to be known
- In many real world problems the environment dynamics is not known before hand => the agent has to *estimate rewards* and *improve its policy* based on direct interaction with the environment

Q-Function

- Instead of a *state value* function, we are more explicit, in storing the value of executing *an action* in a *given state*

$$q^{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a] = E_{\pi}\left[\sum_{\tau=t+1} \gamma^{\tau-t-1} R_{\tau} | S_t = s, A_t = a\right]$$

- Bellman equation for q-function

$$q^{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q^{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a] = \sum_{s' \in S} \sum_{r \in R} p(s', r | s, a) [r + \gamma q^{\pi}(s', \pi(s'))]$$

Q-Learning Algorithm

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **off-policy**

Learning policy is greedy

Play policy allows for exploration

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $random(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{argmax} q(s, a)$   
end
```

```
procedure Q-Learning ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S, a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    while  $s$  not final state do  
      pick action  $a$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma \max_{a'} q(s', a') - q(s, a))$   
       $s \leftarrow s'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \underset{a \in A}{argmax} q(s, a)$   
  end for  
  return  $\pi$ 
```

Q-Learning Algorithm

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **off-policy**

Learning policy is greedy

Play policy allows for exploration

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $random(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{argmax} q(s, a)$   
end
```

```
procedure Q-Learning ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S, a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    while  $s$  not final state do  
      pick action  $a$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma \max_{a'} q(s', a') - q(s, a))$   
       $s \leftarrow s'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \underset{a \text{ in } A}{argmax} q(s, a)$   
  end for  
  return  $\pi$ 
```

Q-Learning Algorithm

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **off-policy**

Learning policy is greedy

Play policy allows for exploration

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $random(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{argmax} q(s, a)$   
end
```

```
procedure Q-Learning ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S, a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    while  $s$  not final state do  
      pick action  $a$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma \max_{a'} q(s', a') - q(s, a))$   
       $s \leftarrow s'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \underset{a \in A}{argmax} q(s, a)$   
  end for  
  return  $\pi$ 
```

Q-Learning Algorithm

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **off-policy**

Learning policy is greedy

Play policy allows for exploration

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $random(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{argmax} q(s, a)$   
end
```

```
procedure Q-Learning ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S, a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    while  $s$  not final state do  
      pick action  $a$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma \underset{a'}{max} q(s', a') - q(s, a))$   
       $s \leftarrow s'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \underset{a \text{ in } A}{argmax} q(s, a)$   
  end for  
  return  $\pi$ 
```

SARSA Algorithm

Agent learns by observing consequences of actions it takes in the environment

Q-values adjusted through **temporal differences**

Learning is **on-policy**

Action used to play/explore =
action used for updating the q-values

```
procedure  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
  with prob  $\epsilon$ : return  $random(A)$   
  with prob  $1-\epsilon$ : return  $\underset{a}{\operatorname{argmax}} q(s, a)$   
end
```

```
procedure SARSA ( $\langle S, A, \gamma \rangle, \epsilon$ )  
  for all  $s$  in  $S, a$  in  $A$  do  
     $q(s, a) \leftarrow 0$  // set initial values to 0  
  end for  
  for all episodes do  
     $s \leftarrow$  initial state  
    pick action  $a$  from  $s$  using  $\epsilon$ -Greedy ( $s, q, \epsilon$ )  
    while  $s$  not final state do  
      execute  $a \rightarrow$  get reward  $r$  and next state  $s'$   
      Pick action  $a'$  from  $s'$  using  $\epsilon$ -Greedy ( $s', q, \epsilon$ )  
       $q(s, a) \leftarrow q(s, a) + \alpha(r + \gamma q(s', a') - q(s, a))$   
       $s \leftarrow s', a \leftarrow a'$   
    end while  
  end for  
  for all  $s$  in  $S$  do  
     $\pi(s) \leftarrow \underset{a \text{ in } A}{\operatorname{argmax}} q(s, a)$   
  end for  
  return  $\pi$ 
```


OpenAI Gym Environments

- Remember the [Taxi-v3](#) and [FrozenLake-v1](#) environment in [OpenAI Gymnasium](#):
- **Task**
 - Implement a Q-Learning agent and a SARSA agent for these environments
 - Create the **reward per training epoch** plot for both **Q-Learning** and **SARSA**
 - Plot the reward evolution for Q-Learning and SARSA **on the same diagram**
 - Every X (e.g. 50, 100) epochs do an **evaluation run** (run the *currently learned policy* for 50 epochs and report the *average reward*)
 - Compare the convergence speed and highest reward metrics of each algorithm under **different hyperparameter settings (see next slide)**

OpenAI Gym Environments

- **Task (details)**
 - **Vary the main parameters influencing learning:**
 - γ – 0.5, 0.9
 - ϵ – 0.1, 0.5, 0.8
 - α (lr) – 0.1, 0.5, 0.9
 - **Analyse the results by:**
 - **Keeping two parameters constant** (e.g. $\gamma=0.9$, $\epsilon = 0.1$) and **varying the third**;
 - **plot all variations of a parameter on the same graph** (e.g. all variations of α are displayed on the same graph to better observe the influence of that single parameter)