

Multi Agent Systems

- Lab 8 -

Deep Q-Network

Q-Learning with Linear Value Function Approximation – General Formulation

- Use *features* to represent state and action $x(s, a) = \begin{pmatrix} x_1(s, a) \\ x_2(s, a) \\ \dots \\ x_n(s, a) \end{pmatrix}$
- Q-function represented as ***weighted linear combination of features***
$$\hat{Q}(s, a, \mathbf{w}) = \mathbf{x}(s, a)^T \mathbf{w} = \sum_{j=1}^n x_j(s, a) w_j$$

- ***Learn*** weights \mathbf{w} through stochastic gradient descent updates
$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \nabla_{\mathbf{w}} E_{\pi}[(Q^{\pi}(s, a) - \hat{Q}^{\pi}(s, a, \mathbf{w}))^2]$$

Q-Learning with Linear Value Function Approximation –

- VFA relies on **features** that have to convey information for learning.
 - Features are **handcrafted** based on human knowledge and intuition about a specific problem
- **Function approximation + off-policy control + bootstrapping** (the “deadly triad”) can fail to converge
- => attempt to improve on one aspect, function approximation, by leveraging Deep Neural Networks as universal function approximators

DQN – TD Target as Loss objective

- For Q-Function, instead of the actual gain per episode under current policy $Q^\pi(s, a)$ use **TD-target** $r + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w})$



- Learn NN** weights \mathbf{w} through stochastic gradient descent updates

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \nabla_{\mathbf{w}} (r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w}))^2 \quad \leftarrow \text{In practice - Huber Loss}$$

$$\Delta \mathbf{w} = \alpha (r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

Experience Replay

- TD Learning Target requires behavioral samples (s, a, r, s')
- Problems in learning Q values through NN trained with SGD arise when inputs (s) are highly correlated (e.g. consecutive frames in a game)
- => alleviate issue through **experience replay**: store a series of behavioral samples (in a **Replay Buffer**), **shuffle them** and **pick a sample** at each training step
- **Replay Buffer** management subject of ongoing research

DQN Procedure

procedure DQN ($\langle S, A, \gamma \rangle, \epsilon, nn_model, mem_size, batch_size$)

 Init **replay_buffer**, Init ***nn_model***, Init ***optimizer***

for all episodes **do**

$s \leftarrow$ initial state

while s not final state **do**

 pick action a using ϵ -Greedy (s, nn_model, ϵ)

 execute $a \rightarrow$ get reward r and next state s'

$replay_buffer.push((s, a, r, s'))$

$learn(nn_model, replay_buffer, batch_size, optimizer)$

$s \leftarrow s'$

end while

end for

procedure learn ($nn, memory, batch_size, optimizer$)

$batch_s, batch_a, batch_r, batch_s' =$
 $memory.sample(batch_size)$

$current_q = nn(batch_s).gather(batch_a)$

$next_state_q = \operatorname{argmax}_a nn(batch_s')$

$td_target = r + \gamma next_state_q$

$loss = \text{huber_loss}(current_q, td_target)$

$optimizer.zero_grad()$

$loss.backward()$

$optimizer.step()$

procedure ϵ -Greedy (s, NN_model, ϵ)

$[\hat{q}(s, a_1), \dots, \hat{q}(s, a_n)] = NN_model(s)$

 with prob ϵ : return $random(A)$

 with prob $1-\epsilon$: return $\operatorname{argmax}_a [\hat{q}(s, a_1), \dots, \hat{q}(s, a_n)]$

end

DQN – Fixed targets

- Standard DQN can still suffer from the “deadly triad” leading to convergence issues
- => improvement: maintain **2 Q-Learning models Q_{model} and Q_{target}** - whereby Q_{target} is “delayed” in updating its parameters => ***fixed target*** (for a while - e.g. 50, 100 steps)
 - Weights of Q_{target} are updated periodically to those of Q_{model}

DQN – Fixed Target Procedure

procedure DQN ($\langle S, A, \gamma \rangle, \epsilon, nn_model, nn_target, mem_size, batch_size, delay_steps$)

Init **replay_buffer**, Init **nn_model**, Init **optimizer**
steps = 0

for all episodes **do**

$s \leftarrow$ initial state

while s not final state **do**

 pick action a using ϵ -Greedy (s, nn_model, ϵ)

 execute $a \rightarrow$ get reward r and next state s'

$replay_buffer.push((s, a, r, s'))$

$learn(nn_model, nn_target, replay_buffer,$
 $batch_size, optimizer)$

$s \leftarrow s'$

 steps += 1

if steps % delay_steps == 0

then $nn_target \leftarrow nn_model$

end while

end for

procedure learn ($model, target, memory, batch_size, optimizer$)

$batch_s, batch_a, batch_r, batch_s' =$

$memory.sample(batch_size)$

$current_q = model(batch_s).gather(batch_a)$

$next_state_q = \argmax_a target(batch_s')$

$td_target = r + \gamma next_state_q$

$loss = \text{huber_loss}(current_q, td_target)$

$optimizer.zero_grad()$

$loss.backward()$

$optimizer.step()$

procedure ϵ -Greedy (s, NN_model, ϵ)

$[\hat{q}(s, a_1), \dots, \hat{q}(s, a_n)] = NN_model(s)$

 with prob ϵ : return $random(A)$

 with prob $1-\epsilon$: return $\argmax_a [\hat{q}(s, a_1), \dots,$
 $\hat{q}(s, a_n)]$

end

Double-DQN

- Standard DQN can still suffer from the “deadly triad” leading to convergence issues
- => improvement: interplay between two modes in Q and Q_{target} in setting the TD target
 - reduce **overestimations** by **decomposing** the max operation in the target into ***action selection*** and ***action evaluation***
 - ***evaluate*** the greedy policy according to the ***online network***, but using the ***target network to estimate its value***.

$$Y_t^{\text{DoubleDQN}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t), \theta_t^-)$$

DDQN Procedure

procedure DQN ($\langle S, A, \gamma \rangle, \epsilon, nn_model, nn_target, mem_size, batch_size, delay_steps$)

Init **replay_buffer**, Init **nn_model**, Init **optimizer**
steps = 0

for all episodes **do**

$s \leftarrow$ initial state

while s not final state **do**

 pick action a using ϵ -Greedy (s, nn_model, ϵ)

 execute $a \rightarrow$ get reward r and next state s'

$replay_buffer.push(s, a, r, s')$

$learn(nn_model, nn_target, replay_buffer,$
 $batch_size, optimizer)$

$s \leftarrow s'$

 steps += 1

if steps % delay_steps == 0

then $nn_target \leftarrow nn_model$

end while

end for

procedure learn ($model, target, memory, batch_size, optimizer$)

$batch_s, batch_a, batch_r, batch_s' =$

$memory.sample(batch_size)$

$current_q = model(batch_s).gather(batch_a)$

$next_state_q =$

$target(batch_s').gather(\argmax_a,$
 $model(batch_s'))$

$td_target = r + \gamma next_state_q$

$loss = \text{huber_loss}(current_q, td_target)$

$optimizer.zero_grad()$

$loss.backward()$

$optimizer.step()$

procedure ϵ -Greedy (s, NN_model, ϵ)

$[\hat{q}(s, a_1), \dots, \hat{q}(s, a_n)] = NN_model(s)$


 with prob ϵ : return $random(A)$

 with prob $1-\epsilon$: return $\argmax_a [\hat{q}(s, a_1), \dots, \hat{q}(s, a_n)]$

end

OpenAI Gym CartPole Environment

- **Cartpole-v1** environment in OpenAI Gymnasium:
 - Objective: keep a pendulum upright for as long as possible
 - 2 actions: left (force = -1), right (force = +1)
 - Reward: +1 for every timestep that the pole remains upright
 - Game ends when pole more the 15° from vertical OR cart moves > 2.4 units from center



OpenAI Gym CartPole Environment

- **Cartpole-v1** environment in **OpenAI Gym**:
 - **Explore with:**
 - The SGD **learning rate**
 - The **target_update_freq** parameter (10, 100, 200)
 - The **learning_freq** parameter (1, 4, 8)
 - **Use** an $\epsilon = \text{decay}(\text{init}=0.9, \text{end}=0.05, \text{nr_iterations})$ – see `eps_generator` in code – **explore different decay rates**
- **Plot** agent learning curves for each case
- Compare agent learning curves for **DQN** against **Double DQN**